

Adaptive model selection for time series prediction in wireless sensor networks

Yann-Aël Le Borgne^{a,*}, Silvia Santini^{b,2}, Gianluca Bontempi^a

^a*ULB Machine Learning Group, Department of Computer Science, Université Libre de Bruxelles (U.L.B.), 1050 Brussels, Belgium*

^b*Department of Computer Science, Institute for Pervasive Computing, ETH Zurich, ETH-Zentrum, IFW D41.2, CH-8092 Zurich, Switzerland*

Received 24 August 2006; received in revised form 28 February 2007; accepted 15 May 2007

Available online 25 May 2007

Abstract

In many practical applications of wireless sensor networks, the sensor nodes are required to report approximations of their readings at regular time intervals. For these applications, it has been shown that time series prediction techniques provide an effective way to reduce the communication effort while guaranteeing user-specified accuracy requirements on collected data. Achievable communication savings offered by time series prediction, however, strongly depend on the type of signal sensed, and in practice an inadequate a priori choice of a prediction model can lead to poor prediction performances. We propose in this paper the adaptive model selection algorithm, a lightweight, online algorithm that allows sensor nodes to autonomously determine a statistically good performing model among a set of candidate models. Experimental results obtained on the basis of 14 real-world sensor time series demonstrate the efficiency and versatility of the proposed framework in improving the communication savings.

© 2007 Elsevier B.V. All rights reserved.

Keywords: Wireless sensor networks; Time series prediction; Prediction models; Adaptive model selection

1. Introduction

In many wireless sensor network deployments, sensor nodes are distributed at various locations over a region of interest and collect data at regular time intervals [1–3]. Each sensor on a node captures

a time series representing the development of the sensed physical variable over space and time. Reporting these time series to a sink³ through the sensor nodes' on-board radio represents a significant communication overhead. Since the radio channel has limited capacity [4] and radio communication is known to be the dominant factor of energy consumption in wireless sensor networks [5,6], the development of adequate data gathering techniques able to reduce the amount of data sent throughout the network is recognized as a key

*Corresponding author. Tel.: +32 2 650 55 94.

E-mail addresses: yleborgn@ulb.ac.be (Y.-A. Le Borgne), santinis@inf.ethz.ch (S. Santini), gbonte@ulb.ac.be (G. Bontempi).

¹Supported by European Community COMP2SYS project (MEST-CT-2004-505079).

²Partially supported by the European Community Embedded WiSeNts project (FP6-004400) and by the Swiss National Science Foundation (NCCR-MICS, grant number 5005-67322).

³We refer to a *sink node* as either the central server or as a sensor node responsible to relay data further to the central server.

factor for allowing long-term, unattended network operation. This topic has therefore gained increasing attention in the wireless sensor network research community [7–14].

In many practical application scenarios for sensor networks, users are interested in observing physical phenomena with a pre-specified, application-dependent accuracy. Thus, gathered sensor data are usually accepted to lie within a known error bound, say $[-\varepsilon, +\varepsilon]$, $\varepsilon \in \mathbb{R}^+$. A sensor node regularly collecting local measurements can fit a prediction model to the real data and communicate it to the sink, which can then use the model to compute estimates of future sensor readings [8]. Since the sensor node can reproduce the same readings estimations, it transmits a model update to the sink only if the current measurement differs from the predicted by more than $\pm\varepsilon$, thus avoiding unnecessary communication. In the absence of notification from the sensor node, the sink implicitly assumes that the value obtained from the shared prediction model is within the required error bound. This strategy, which we will refer to as *dual prediction scheme* or *DPS* henceforth,⁴ has been shown to lead to potentially high communication and energy savings once adequate prediction models are used [8,9,11,13,14]. Typically, the model to use (e.g. constant, linear, etc.) is fixed a-priori, while model parameters are estimated on the basis of incoming data [11,13,14].

Fig. 1 illustrates how the DPS behaves on a temperature time series obtained from a real-world sensor deployment [16], when the required data accuracy ε is set to 0.5°C and a simple autoregressive (AR) model is used. We can observe that the predicted data are within $\pm 0.5^\circ\text{C}$ of the real data up to the 1261st time step. At time $t = 1262$, the prediction error exceeds the tolerated threshold ε and the sample collected at time $t = 1262$ is sent to the sink. The prediction model is then updated to take into account the new acquired data and from time $t = 1263$ to $t = 1272$, the predicted measurements are again close enough to the real ones, making the communication between the sensing node and the sink superfluous. At $t = 1273$, the sensor node realizes again that the sink is predicting the sensor measurements with an error bigger than ε and thus transmits the current reading. The procedure is repeated again at $t = 1286$, when a




Fig. 1. Dual prediction scheme with error threshold $\varepsilon_{\max} = 0.5^\circ\text{C}$.

new update is sent to the sink. In this example, out of 35 collected readings, only 3 were effectively transmitted by the sensor node, which amounted to about 90% of communication savings. Obviously, the achievable communication savings depend, among others, on the particular sensed phenomenon, on the data sampling rate and, last but not least, on the used prediction model.

Different approaches have been proposed to perform time series forecasting in wireless sensor networks [11,13,14], ranging from simple heuristics to sophisticated modelling frameworks. All these methods allowed to improve upon the simple monitoring approach in which measurements are continuously reported at fixed time intervals. However, these implementations overlooked two relevant issues. First, complex prediction techniques, like, e.g. Kalman filtering, rely on parameters whose identification proves to be difficult in practical settings, particularly when no a-priori knowledge on the signals is available. These difficulties increase with the flexibility of the model, or, equivalently, with the number of parameters necessary to specify it. Therefore, the more flexible the model, the less usable in practice. Second, as the DPS requires the sensor node and the sink to run the same prediction model, all the parameters of the model must be sent every time an update is needed. There exists in terms of communication savings, a tradeoff between the ability of a model to properly fit the signal so as to lower the number of data transmissions and model updates, and the number of parameters that need to be sent at each update.

⁴Other authors dubbed this strategy *dual prediction reporting* (DPR) [15].

In this paper, we address both these issues and we propose a generic procedure for adaptive model selection, henceforth referred to as AMS. The rationale of our approach is to use complex prediction models only if they prove to be efficient both in terms of computation and achievable communication savings, and otherwise to rely on simpler models. To this end, a set of models of increasing complexity is considered, and their performances are assessed in an online fashion, as sensor data are collected, on the basis of a metric that weights the number of updates by their size. In this way, it is possible to select, among a set of candidates, the model that offers at each time instant the highest achievable communication savings. Supported by the literature on time series prediction [17,18], we propose an implementation of the AMS based on AR models, whose parameters can be updated in an online fashion as new observations become available, and that are computationally thrifty to maintain. We show on the basis of 14 publicly available time series captured by real-world sensor nodes that gains achievable by complex prediction models quickly drop as the number of parameters increase, and that, therefore, very few models effectively need to be considered in practical settings. Finally, we propose to rely on a statistical procedure known as *racing* [19], to discard over time models that perform poorly so as to save sensor nodes' computational and memory resources. The obtained experimental results, presented in detail in Section 5, show that the AMS provides a lightweight and efficient implementation of the DPS.

The remainder of this paper is organized as follows: Section 2 introduces the general framework of the DPS and the limits of previously proposed approaches. Section 3 describes the proposed AMS algorithm. Experimental setup and results are reported in Sections 4 and 5, respectively. Finally, limits and potential improvements are discussed in Section 6.

2. The dual prediction scheme

Continuously reporting sensor readings to a sink node at regular time intervals is the most widely used data gathering mechanism in real wireless sensor networks deployments [1,20]. For this reason, we refer to it as the *default monitoring* data collection scheme. With respect to the default monitoring scheme, the DPS significantly reduces

communication between a sensor node and the sink, while guaranteeing the data collected to be within a user-specified accuracy. The gains in communication offered by the DPS, however, depends on the adequacy between the prediction model used and the characteristics of the time series captured by the sensor. Providing an overview on time series forecasting techniques is beyond the scope of this paper, and an interested reader may refer to [18,21–23]. In this section, we outline the particular problems and challenges that arise when applying the DPS in wireless sensor networks.

2.1. Time series prediction models

Let $\mathcal{X}_t = \langle X_0, X_1, X_2, \dots \rangle$ be a time series representing the sequence of collected sensor measurements $X_t \in \mathbb{R}$ for each sampling time instant $t \in \mathbb{N}$, and let \hat{X}_t indicate an estimation of the element X_t at time t . Let $\mathbf{X}_{[0:t]} = \langle X_0, X_1, X_2, \dots, X_{t-1}, X_t \rangle$ be the sequence of observations up to time t . Let

$$\hat{X}_{t+1} = h(\mathbf{X}_{h,t}, \boldsymbol{\theta}_{h,t}) \quad (1)$$

be the estimate of X_{t+1} returned by a prediction model $h(\mathbf{X}_{h,t}, \boldsymbol{\theta}_{h,t})$ on the basis of $\mathbf{X}_{[0:t]}$. A prediction model $h(\mathbf{X}_{h,t}, \boldsymbol{\theta}_{h,t})$ for the time series \mathcal{X}_t is a mapping that takes as inputs a row vector of input values $\mathbf{X}_{h,t}$ (a subset of $\mathbf{X}_{[0:t]}$), together with a row vector of parameters $\boldsymbol{\theta}_{h,t} = (\theta_1, \theta_2, \dots, \theta_k)$, with $k \in \mathbb{N}^+$, and that returns an estimate \hat{X}_{t+1} . We shall in the following refer to $h(\mathbf{X}_{h,t}, \boldsymbol{\theta}_{h,t})$ as h for short. Vectors $\mathbf{X}_{h,t}$ and $\boldsymbol{\theta}_{h,t}$ depend on the model h . For example, a constant model requires just one input value and no parameters, while an AR model of order 2 takes 2 input values and 2 parameters (cf. Section 4.1). These vectors also depend on t if input values and parameters vary over time (see for example Section 4.1 for details). These dependencies are therefore expressed for input values $\mathbf{X}_{h,t}$ and parameters $\boldsymbol{\theta}_{h,t}$ by the means of the subscript (h, t) .

A variety of different prediction models can be used to perform time series forecasting [17,18,23]. The choice of an adequate model depends on the nature of the time series, on the amount of available a-priori knowledge, on the required forecasting accuracy, as well as on the available computational resources.

2.2. DPS: overview and limits

The main task of the DPS is to run an *identical* prediction model h at both the source and the sink

nodes and to use it to produce estimates of the future sensor readings, given some of the previous samples. If the predicted value differs from the actual sensor measurements by more than a given error threshold ε ,⁵ a model update⁶ is transmitted to the sink.

The simplest implementation of the DPS uses a constant prediction model, henceforth referred to as CM, which allows the sink to reconstruct a piecewise constant approximation of the real sensor signal. Using a CM, no updates are sent as long as readings collected by the sensor do not diverge by more than $\pm\varepsilon$ from the last reading sent to the sink. When this difference becomes bigger than ε , the current reading is sent, and this process is repeated over time. This approach, proposed by Lazaridis and Mehrotra [9] and Olston et al. [8], provides appealing communication savings with respect to the default monitoring strategy. On selected time series, a CM can be outperformed by more complex prediction techniques, as shown in [11,13,14]. However, all these methods depend on a number of parameters that cannot be fixed a priori in a generic manner.

An example is the instantiation of the DPS using Kalman filters [11]. Kalman filters are based on a sound theoretical framework, and are particularly attractive for implementing the DPS as they allow to model a large variety of phenomena. However, their use requires to define a priori a *state transition matrix*, which describes the underlying dynamics of the signal, and two covariance matrices specifying the process noise and the observation noise. In [11], authors for example had to arbitrarily specify the process and observation noise to be diagonal matrices with value 0.05. As in most cases such information cannot be set a priori, simpler approaches have been studied. Santini and Römer propose the adoption of autoregressive adaptive filters, updated by the means of the LMS procedure, which are able to learn signal statistics on the fly, and can continuously and autonomously adapt to changes [14]. At the same time, Tulone and Madden [13] interestingly cast the DPS in a more general

framework, applying AR models and including the identification of outliers, whose appearance is not reported to the sink. In both cases, however, neither online procedures for setting the order of the AR model nor considerations about the communication overhead caused by the transmission of the necessary model parameters are provided. As we will show in Section 5, these issues are of critical importance as they can seriously impact the achievable communication savings.

3. Adaptive model selection

This section presents the AMS strategy that we propose for allowing a sensor node to autonomously select a model out of a set of possible prediction models. A sensor node running AMS maintains a set of K *candidate* prediction models $h_i(\mathbf{X}_{h_i,t}, \boldsymbol{\theta}_{h_i,t})$, $1 \leq i \leq K$. For each model h_i , a given quality measure is recursively estimated and the model that optimizes this performance indicator is selected as the *current* model. These performance measures, presented in Section 3.1, may also be used to run the *racing* mechanism, which allows to discard poorly performing models from the set of candidate models, as we describe in Section 3.2. The detailed AMS algorithm is presented in Section 3.3.

3.1. Performance estimates

The main goal of the DPS is to reduce the number of updates between a sensor node and the sink. To measure the performance of the DPS it is therefore meaningful to consider the *relative update rate*, i.e. the ratio of the number of updates effectively sent when running the DPS to the number of updates that would have been sent by the default monitoring scheme. Let $U_{h_i,t}$ be the relative update rate for the model h_i at time t , where $U_{h_i,1} = 1$. $U_{h_i,t}$ can be recursively computed as $U_{h_i,t} = \frac{(t-1)U_{h_i,t-1} + 1}{t}$ if an update is needed at time t , or as $U_{h_i,t} = \frac{(t-1)U_{h_i,t-1}}{t}$ otherwise. The relative update rate reflects the percentage of transmitted packets with respect to the default monitoring scheme. Note that the relative update rate for the default monitoring scheme is 1 since it requires the transmission of the totality of the readings, and that any lower value indicates a gain in the number of transmitted packets.

Performance assessment in terms of update rate has been considered in several implementations of

⁵To work properly, the DPS requires all sent data to actually reach the sink, thus a loss-free or acknowledged communication link between the node and the sink is required. See [14] and Section 6 for a discussion of some mechanisms which relax this assumption.

⁶A model update can consist in either a variation in input values and model parameters or in the choice of a brand-new model.

the DPS [9,11,13,14]. However, this performance indicator does not take into account the fact that while an update in the default monitoring mode only consists of the current sensor readings, updating a model h_i requires the input values $\mathbf{X}_{h_i,t}$ and the model parameters $\theta_{h_i,t}$. Consequently, performing a single update may require sending a high number of bytes to the sink, which may become critical in settings characterized by a very limited network bandwidth. To take into account the packet size of a single model update we introduce an alternative performance indicator, *relative data rate*, which we define as follows:

$$W_{h_i,t} = U_{h_i,t} \times C_{h_i}, \quad (2)$$

where C_{h_i} (henceforth referred to as *model cost*) is the ratio of the number of bytes required to send an update of model h_i to the number of bytes required to send an update in the default monitoring mode. The relative data rate $W_{h_i,t}$ measures the savings in terms of data rate for model h_i at time t with respect to the default monitoring mode.

3.2. Racing mechanism

Since it is likely that some $\{h_i\}$ will perform poorly, it would be preferable not to maintain them in order to save computational and memory requirements. An effective approach to detecting prediction models that perform poorly out of a set of candidate models is offered by the *racing* mechanism [19]. The rationale of the racing mechanism is to determine, on the basis of hypothesis testing [24], what models among a set of candidate models are significantly outperformed by others. For instance, let $h_i^* = \operatorname{argmin}_{h_i} W_{h_i,t}$ be the model with the lowest relative data rate at time instant t among the set of candidate models $\{h_i\}$, and let $\Delta_{h_i,h_i^*} = W_{h_i,t} - W_{h_i^*,t}$ be the difference between the estimated relative data rates of any model h_i and h_i^* . Relying on the Hoeffding bound [25], a distribution-free statistical bound, the racing mechanism assumes with probability $1 - \delta$ that h_i^* truly outperforms h_i if

$$\Delta_{h_i,h_i^*} > R \sqrt{\frac{\ln(1/\delta)}{2t}}, \quad (3)$$

where R is the range taken by the random variable Δ_{h_i,h_i^*} . Thanks to the lack of parametric assumptions, the Hoeffding bound requires no other information than the range of values taken by the random variables considered, which is known in

Table 1
Adaptive model selection algorithm

Adaptive model selection algorithm	Algorithm for virtual model updates
Algorithm AMS($K, \{h_i\}, \{C_{h_i}\}, \varepsilon$)	Algorithm simulateModel(h_i, X_t)
$U_{h_i,1} \leftarrow 1$ for $1 \leq i \leq K$	$\hat{X}_{t+1} \leftarrow \text{predictValue}(h_i^*)$
$h^* \leftarrow \operatorname{argmin}_{h_i} C_{h_i}$	update(h_i, X_{t+1})
While True	if $(\hat{X}_t - X_t > \varepsilon)$
$X_t \leftarrow \text{getNewReading}()$	$U_{h_i,t} \leftarrow \frac{(t-1) * U_{h_i,t-1} + 1}{t}$
For (i in $1 : K$)	$\theta_{h_i,t}^* \leftarrow \theta_{h_i,t}$
$h_i \leftarrow$	else
simulateModel(h_i, X_t)	$U_{h_i,t} \leftarrow \frac{(t-1) * U_{h_i,t-1}}{t}$
endFor	endif
$\{h_i\} \leftarrow \text{racing}(\{h_i\})$ see	
Equation (3)	Return h_i
$\hat{X}_t \leftarrow \text{predictValue}(h^*)$	
if $(\hat{X}_t - X_t > \varepsilon)$	
$h^* \leftarrow \operatorname{argmin}_{h_i} U_{h_i,t} * C_{h_i}$	
sendNewModel(h^*)	
endif	
endWhile	

advance. As $0 \leq W_{h_i,t} \leq C_{h_i}$ and $0 \leq W_{h_i^*,t} \leq C_{h_i^*}$, it follows $R = C_{h_i} + C_{h_i^*}$, and the bound for discarding model h_i is therefore given by

$$\Delta_{h_i,h_i^*} > (C_{h_i} + C_{h_i^*}) \sqrt{\frac{\ln(1/\delta)}{2t}}. \quad (4)$$

The racing mechanism discards poor performing models from the set of candidates among which the AMS chooses the current model. Since the bound gets tighter as t increases, only one model is eventually maintained on the sensor node.

3.3. AMS algorithm

Table 1 contains a pseudo-code version of the AMS algorithm. It takes as inputs the error tolerance ε , the number of candidate models K , the set of models $\{h_i\}$, and their corresponding costs $\{C_{h_i}\}$.⁷ The first model sent to the sink is that with the lowest model cost. When the sensor collects a new reading X_t , the AMS runs the function *simulateModel*, which estimates the relative update rates $U_{h_i,t}$ for all candidate models h_i . This function first determines whether an update is necessary or not by checking if the current reading estimation $\hat{X}_t = h_i(\mathbf{X}_{h_i,t-1}, \theta_{h_i,t-1}^*)$, computed by model h_i at

⁷The model costs must all be set to 1 if the relative update rate is used as performance indicator

time t , is more than $\pm\epsilon$ off the actual sensor value X_t . The relative update rate $U_{h_i,t}$ is then computed as described in Section 3.1. Moreover, since the parameters of a candidate model may be updated recursively as new sensor readings become available, the function *simulateModel* maintains two sets of parameters for each model h_i : $\theta_{h_i,t}$ and $\theta_{h_i,t}^*$. Parameters $\theta_{h_i,t}$ are continuously updated with incoming data so that the model is constantly refined (e.g. using the recursive least square (RLS) procedure for AR models, as detailed in Section 4.1). On the contrary, as long as no update is necessary for model h_i , parameters $\theta_{h_i,t}^*$ remain unchanged since they represent the parameters that would be shared by the sensor node with the sink if h_i was the current model.

After running to completion, the function *simulateModel* returns the control to AMS, which then behaves as it were a “classical” DPS scheme. It therefore checks whether the absolute value of the difference between the reading estimation $\hat{X}_{t+1} = h^*(\mathbf{X}_{h^*,t}, \theta_{h^*,t})$, computed at the sink using the current model h^* , and the actual sensor value X_t does not exceed the tolerated error threshold ϵ . If this threshold is exceeded, the current model h^* is assigned the model in $\{h_i\}$ that minimizes the chosen performance indicator, and an update composed of the input values $\mathbf{X}_{h^*,t}$ and the parameters $\theta_{h^*,t}$ is sent to the sink.

4. Experimental setup

In this section, we describe the setup we used to assess the performance of the AMS algorithm. The corresponding experimental results are reported in Section 5.

4.1. Prediction models

We based the experimental evaluation of our AMS algorithm on AR models. We tested how AR models, whose parameters can be recursively updated, can improve upon a CM when running the DPS. AR models have been chosen for two reasons. First, they have been shown to be both theoretically and experimentally good candidates for time series predictions [17,18]. Second, model parameters can be estimated by the means of the RLS algorithm [26], which allows to adapt the parameters to the underlying time series in an online fashion, without the need of storing large sets of past data.

Time series forecasting using AR models is performed by regressing the value X_t of the time series \mathcal{X}_t at time instant t against the elements of the time series at the previous p time instants $(X_{t-1}, X_{t-2}, \dots, X_{t-p})$. The prediction at time $t+1$ is thus obtained as

$$\hat{X}_{t+1} = \theta_1 X_t + \theta_2 X_{t-1} + \dots + \theta_p X_{t-p+1}, \quad (5)$$

where $(\theta_1, \theta_2, \dots, \theta_p)$ are the AR coefficients and p is the *order* of the AR model, thus denoted as $\text{AR}(p)$. Following the notations introduced in Section 3, let $\theta_{\text{AR}(p),t} = (\theta_1(t), \theta_2(t), \dots, \theta_p(t))$ be the row vector of parameters and $\mathbf{X}_{\text{AR}(p),t} = (X_t, X_{t-1}, \dots, X_{t-p+1})$ be the row vector of inputs to a model $\text{AR}(p)$. The prediction is thus obtained by the scalar product:

$$\hat{X}_{t+1} = \theta_{\text{AR}(p),t} \cdot \mathbf{X}_{\text{AR}(p),t}^T, \quad (6)$$

where the superscript T stands for the transposition operator. The parameters $\theta_{\text{AR}(p),t}$ can be computed by means of the RLS algorithm, which consists in a computationally thrifty set of equations that allows to recursively update the parameters $\theta_{\text{AR}(p),t}$ as new observations X_t are available. The computational cost for an update of the vector $\theta_{\text{AR}(p),t}$ is of $3p^3 + 5p^2 + 4p$.

4.2. Datasets

The experimental evaluation is based on a set of 14 publicly available datasets, collected in real sensor network deployments. The datasets vary in terms of the nature of the observed phenomenon, signal dynamic, sampling frequency and length, and are briefly listed in Table 2.

4.3. Generic error threshold for performance comparison

To be able to compare results obtained from different datasets regardless of the specific physical quantities being examined, the influence of the threshold parameter ϵ is analyzed by considering it as proportional, through a given factor k , to the range r of the signal. The range r was computed by taking the difference between the maximal and minimal values in the time series. The case $k = 0.01$ accounts for scenarios in which high precision is required, while $k = 0.2$ corresponds to a very rough bound on the tolerated error.

Table 2
Data sets

Data set	Sensed quantity	Sampling period	Period	Number of samples	Source
S Heater	Temperature	3 s	–	3000	[27]
I Light	Light	5 min	8 days	1584	[20]
M Hum	Humidity	10 min	30 days	4320	[16]
M Temp	Temperature	10 min	30 days	4320	[16]
NDBC WD	Wind direction	1 h	1 year	7564	[28]
NDBC WSPD	Wind speed	1 h	1 year	7564	[28]
NDBC DPD	Dominant wave period	1 h	1 year	7562	[28]
NDBC AVP	Average wave period	1 h	1 year	8639	[28]
NDBC BAR	Air pressure	1 h	1 year	8639	[28]
NDBC ATMP	Air temperature	1 h	1 year	8639	[28]
NDBC WTMP	Water temperature	1 h	1 year	8734	[28]
NDBC DEWP	Dewpoint temperature	1 h	1 year	8734	[28]
NDBC GST	Gust speed	1 h	1 year	8710	[28]
NDBC WVHT	Wave height	1 h	1 year	8723	[28]

5. Experimental results

In this section, we report extensive experimental results to assess the performance of the AMS algorithm. First, we report the communication gains achievable running the “classical” DPS with the CM and with AR models of orders 1–5 (AR1, ..., AR5), both in terms of relative update rate (in Section 5.1) and in terms of relative data rate (in Section 5.2). Along with these results we highlight the benefits of using the AMS, which for all time series was able to select the best model. Results for the convergence rate of the racing mechanism are reported in Section 5.3, and average gains in data rate obtained on all 14 time series as a function of the tolerated error threshold are presented in Section 5.4.

5.1. Gains in update rate

Table 3 reports the percentage of packets sent when running the DPS with the CM and AR models with orders from 1 to 5 (AR1, ..., AR5). The error tolerance was fixed at $0.01 \times r$, and results are reported for each of the 14 time series presented in Section 4.2. Bold-faced figures indicate models that are not significantly outperformed by the model with the lowest update rate (one tailed t -test with respect to best model, $p < 0.05$).

We remark that in most cases, AR models outperformed the CM, and that performances of AR models are statistically equivalent regardless of the model order. However, the CM performed significantly better than AR models for three time

Table 3

Percentage of transmitted packets for DPS run with different time series forecasting methods

	CM	AR1	AR2	AR3	AR4	AR5	AMS
S Heater	74	75	61	59	59	59	AR3
I Light	38	40	39	40	40	39	CM
M Hum	53	53	49	50	49	49	AR4
M Temp	48	48	45	45	44	44	AR4
NDBC DPD	65	85	80	80	80	80	CM
NDBC AWP	72	73	73	73	73	73	CM
NDBC BAR	51	50	39	39	39	37	AR5
NDBC ATMP	39	39	36	36	36	36	AR3
NDBC WTMP	27	27	21	21	21	20	AR5
NDBC DEWP	57	52	52	52	52	52	AR3
NDBC WSPD	74	84	82	83	83	83	CM
NDBC WD	85	81	81	81	81	81	AR1
NDBC GST	80	81	80	80	80	81	CM
NDBC WVHT	58	56	56	56	56	56	AR3

Bold-faced numbers indicate models that yielded the best performances (one tailed t -test with respect to best model, $p < 0.05$).

series, namely I Light, NDBC DPD and NDBC WSPD, and yielded similar performances for NDBC AWP and NDBC GST. These apparent deficiencies of AR models are due to the nature of those time series, qualitatively characterized by sudden and sharp changes. These abrupt changes cause the variance in the estimation of AR coefficients to increase, making the models unstable and thus allowing a simple CM to provide better performances in terms of update rates (with gains of about 15% with respect to AR models for NDBC DPD and gains up to 8% for NDBC WSPD over a 1 year period).

The last column of Table 3 contains the model that yielded the lowest update rate, and that was consequently selected by the AMS procedure.

5.2. Gains in data rate

In this section we assess the performances of the DPS in terms of the weighted update rate $W_{h_i,t} = U_{h_i,t} \times C_{h_i}$, or data rate, introduced in Section 3.1. Model costs C_{h_i} were computed assuming that each data sample and parameter can be stored in one byte. Accordingly, the CM requires 1 byte to be sent to the sink, while the update of an AR(p) model requires $2p$ bytes (i.e. p bytes of data as initial input values for the model, along with p bytes of data for sending the parameters). The packet overhead (header and footer) depends on the specific communication protocol. For our experiments, we considered a packet overhead of $P_{\text{overhead}} = 24$ bytes, which corresponds to the average overhead of the IEEE 802.15.4 protocol, a potential standard for wireless sensor networks [6]. The size of a packet carrying an update for an AR(p) model is therefore

$$C_{\text{AR}(p)} = \frac{24 + 2p}{24 + 1}. \quad (7)$$

Table 4 reports the performances of the CM and AR(p) models in terms of percentage of bytes sent to the sink with respect to the default monitoring mode. Note that as the cost of the CM is 1, figures

Table 4
Percentage of transmitted bytes for DPS run with different time series forecasting methods

	CM	AR1	AR2	AR3	AR4	AR5	AMS
S Heater	74	78	68	70	76	81	AR2
I Light	38	42	44	48	51	53	CM
M Hum	53	55	55	60	62	66	CM
M Temp	48	50	50	54	56	60	CM
NDBC DPD	65	89	89	95	102	109	CM
NDBC AWP	72	75	81	88	93	99	CM
NDBC BAR	51	52	44	47	49	50	AR2
NDBC ATMP	39	41	40	43	46	49	CM
NDBC WTMP	27	28	23	25	27	28	AR2
NDBC DEWP	57	54	58	62	67	71	AR1
NDBC WSPD	74	87	92	99	106	113	CM
NDBC WD	85	84	91	98	104	111	AR1
NDBC GST	80	84	90	96	103	110	CM
NDBC WVHT	58	58	63	67	71	76	CM

Bold-faced numbers indicate models that yielded the best performances (one tailed t -test with respect to best model, $p < 0.05$).

of the first column of Tables 3 and 4 are identical. In contrast, there is a general deterioration of performances of AR models, as the cost associated with sending their parameters lead them to lose their advantage in terms of prediction accuracy over more simple models. Out of all tested time series, AR models only outperformed the CM five times (on *S Heater*, *NDBC BAR*, *NDBC WTMP*, *NDBC DEWP*), and models eventually selected by AMS were AR(2) (three times) and AR(1) (twice). The AMS column contains the model that yielded the lowest data rate for each time series.

5.3. Racing mechanism

We report in this section the convergence speed obtained when relying on the racing mechanism. Fig. 2 shows the average number of models, averaged over the 14 time series and the first 1000 time instants. The weighted update rate $W_{h_i,t}$ was used to evaluate performance of competing models (as in Section 5.2), with a confidence $1 - \delta = 0.95\%$. Efficiency of the racing in terms of rapidity in discarding poorly performing models depends on the nature of the time series. The convergence to the best model in less than 1000 time instants was obtained in four cases. For other cases, subsets of 2 or 3 remaining models were still in competition after 1000 time instants. The performances of remaining models were in those cases ranging from less than 1% up to 5%, and the a posteriori best model was always observed to be part of the remaining set. AR(4) and AR(5) were discarded in all cases due to




Fig. 2. Number of remaining models over time.

the overhead incurred by sending their parameters to the sink. For five time series, AR(3) and AR(4) were in the remaining candidate models, while for the other nine time series, either CM, AR(1), or both were still competing after the 1000th time instant.

5.4. Tolerated prediction error

This section presents the gains in data rate obtained as the accuracy threshold ε is relaxed. Fig. 3 reports the percentage of transmitted bytes for each of the 14 time series as the error tolerance is relaxed. AMS was run using the relative data rate of competing models as performance indicator. Interestingly, for a $0.05 \times r$ accuracy threshold, which corresponds to good approximation of the sensed phenomenon, less than 20% of data were sent to the sink in comparison to the default monitoring scheme. This reduction decreased down to only 5% of bytes transmitted for an error tolerance of $0.2 \times r$. We should also notice that as the error tolerance increases, the predictive capacity of any method tends to converge to that of the CM. Thus, for error tolerance bigger than $0.1 \times r$, the AMS does not perform, in general, significantly better than the CM (results not reported for space constraints).

5.5. Summary

Applying the DPS with a CM yielded significant communication savings with respect to the default




Fig. 3. Percentage of packet transmitted as tolerance on error prediction is relaxed (proportional to the range r).

monitoring mode, ranging from about 50%, for a very tight error threshold (one hundredth of the sensor signal range), to 96% for rough approximations (one fifth of the sensor signal range). The communication savings could be further slightly reduced on average by relying on AR models, although only those with low orders (less than three with the estimates proposed in Section 5.2) were observed to provide improvements due to the additional communication overhead incurred by transmitting their parameters. Out of the set of initial candidate models, the application of the AMS procedure allowed to select in an online manner the best performing one, which turned out to be the CM for some of the signals.

An empirical guideline that therefore seems to be adapted to scenarios where no a priori knowledge is available on the signal characteristics is to apply the AMS procedure with a set of about four models, composed of the CM and AR models up to order three. Aspects related to the additional computational cost incurred by the AMS procedure, and to alternatives to AR models are addressed in the following section.

6. Discussion

The previous section assessed the ability of the AMS strategy to improve the DPS. We here wish to discuss general issues concerning the DPS.

- (1) *Initial set of candidate models:* We particularly focused in this paper on monitoring scenarios where no a priori knowledge is available on the signal, and motivated the use of AR models on the basis of their popularity and their low computational and memory costs. The proposed AMS procedure can, however, be readily applied to monitoring scenarios where some a priori knowledge is available, by relying on more specific modelling techniques offered in the time series prediction literature. For example, PARIMA models could be used if the observed signal is known to be periodic, or Kalman filters could provide an adequate framework if the signal underlying dynamics are known.
- (2) *Computational overhead:* Although the AMS aims at improving the reduction of the communication between the network nodes, it is also desirable that its computational cost be kept low enough to avoid an excessive energy

consumption. We provide in the following a rough comparison in order to suggest that the computational cost can be deemed negligible for running the AMS with the guidelines proposed in Section 5.5. Running a CM and AR models up to order 3 would require to perform about 200 computations^{8,9} at each time step (see Section 4.1). Given that the ratio of the energy spent in sending one bit of information to the energy spent in executing one instruction has been estimated to be around 2000 for a variety of real sensor network platforms [29], sending a packet of 26 bytes (208 bits) therefore amounts to, in terms of energy consumption, 416 000 CPU instructions, i.e. running the AMS over about 2000 time instants. This rough estimate suggests that the energy cost related to running the AMS may in most cases be compensated by the energy gained with the AMS communication savings.

- (3) *Network unreliability*: The main concern raised by relying on a DPS resides in the fact that in absence of notification from a sensor node, the sink deems the prediction given by the shared model to fall within the ε error tolerance. Additional checking procedures must therefore be considered for the DPS to be reliable, as an absence of notification can also derive from packet loss (as wireless links are particularly lossy) or sensor failure. To that end, a “watchdog” regularly checking the sensor activity and the number of sent packets can be set up, as already discussed in [14]. By keeping in the sensor node memory the recent history of sent updates, these can be communicated to the sink at checking intervals if the number of packets sent differs from the number of packets received. Node failure is detected by the absence of acknowledgment from the sensor node to the watchdog request. Choice for the watchdog period depends on the application, as critical application such as fire detection would require low latency in being warned that a sensor node fails to communicate, whereas for other applica-

tions, i.e. field watering or air conditioning system, that period could be made larger.

- (4) *Outlier detection*: Outliers, i.e. erroneous readings, can appear in the flow of collected readings. Detecting outliers on the sensor node is an important issue if those are frequent, as they both entail useless update packets, and jeopardize the convergence of model parameters. Detection of outliers is, however, a nontrivial task [30]. We mention briefly as solutions to this issue the possibility of adding statistical tests on measured readings as discussed in [13], or a priori bounds on measured readings (e.g. all temperature readings outside the range $[-20^{\circ}\text{C}; +50^{\circ}\text{C}]$ should be ignored). Note that using prediction models allows in cases of detected outliers a way to replace the erroneous value with a more likely value.

7. Conclusions

In this paper, we proposed the AMS, a generic algorithm for online model selection of time series prediction models for reducing the communication effort in wireless sensor networks. Our adaptive selection scheme aimed at making sensor nodes smart enough to be able to autonomously and adaptively determine the best performing model to use for performing data prediction. In order to develop a general applicable framework, we proposed a possible implementation relying on CM and AR models that works in a fully online manner. Extensive evaluation on a set of 14 time series demonstrated the ability of our algorithm to significantly reduce the number of data transmissions for all time series while complying with the poor available memory and computational resources of common sensor network platforms.

References

- [1] R. Szwedczyk, A. Mainwaring, J. Polastre, D. Culler, An analysis of a large scale habitat monitoring application, in: Proceedings of the 2nd ACM Conference on Embedded Networked Sensor Systems (SenSys'04), Baltimore, MD, USA, 2004.
- [2] M.A. Batalin, M.H. Rahimi, Y. Yu, D. Liu, A. Kansal, G.S. Sukhatme, W.J. Kaiser, M. Hansen, G.J. Pottie, M.B. Srivastava, D. Estrin, Call and response: experiments in sampling the environment, in: Proceedings of the 2nd ACM Conference on Embedded Networked Sensor Systems (SenSys'04), Baltimore, MD, USA, 2004, pp. 25–38.

⁸The actual number of CPU instructions corresponding to an addition or multiplication may vary depending on the particular microprocessor. There is no loss in generality in considering here the number of additions and multiplications to be equal to the number of CPU instructions.

⁹The number of computations was obtained by taking the computational cost of the RLS algorithm for AR models with orders 1–3.

- [3] P. Buonadonna, D. Gay, J.M. Hellerstein, W. Hong, S. Madden, TASK: sensor network in a box, in: Proceedings of the 2nd IEEE European Workshop on Wireless Sensor Networks and Applications (EWSN'05), Istanbul, Turkey, 2005.
- [4] D. Marco, E.J. Duarte-Melo, M. Liu, D.L. Neuhoff, On the many-to-one transport capacity of a dense wireless sensor network and the compressibility of its data, in: 2nd International Workshop on Information Processing in Sensor Networks (IPSN'03), Springer, Palo Alto, CA, USA, 2003.
- [5] I.F. Akyildiz, W. Su, Y. Sankarasubramaniam, E. Cayirci, Wireless sensor networks: a survey, *Comput. Networks* 38 (4) (2002) 393–422.
- [6] J. Polastre, R. Szewczyk, D. Culler, Telos: enabling ultra-low power wireless research, in: Proceedings of the 4th International Conference on Information Processing in Sensor Networks: Special Track on Platform Tools and Design Methods for Network Embedded Sensors (IPSN'05/SPOTS), 2005.
- [7] S. Goel, T. Imielinski, Prediction-based monitoring in sensor networks: taking lessons from MPEG, *SIGCOMM Comput. Commun. Rev.* 31 (5) (2001) 82–98.
- [8] C. Olston, J. Jiang, J. Widom, Adaptive filters for continuous queries over distributed data streams, in: Proceedings of the 2003 ACM SIGMOD International Conference on Management of Data (SIGMOD'03), ACM Press, New York, NY, USA, 2003, pp. 563–574.
- [9] I. Lazaridis, S. Mehrotra, Capturing sensor-generated time series with quality guarantee, in: Proceedings of the 19th International Conference on Data Engineering (ICDE'03), Bangalore, India, 2003.
- [10] A. Deshpande, C. Guestrin, S. Madden, J. Hellerstein, W. Hong, Model-driven data acquisition in sensor networks, in: Proceedings of the 30th Very Large Data Base Conference (VLDB'04), Toronto, Canada, 2004.
- [11] A. Jain, E.Y. Chang, Y.-F. Wang, Adaptive stream resource management using kalman filters, in: Proceedings of the ACM SIGMOD International Conference on Management of Data (SIGMOD'04), ACM Press, New York, NY, USA, 2004.
- [12] Y. Le Borgne, G. Bontempi, Round robin cycle for predictions in wireless sensor networks, in: 2nd International Conference on Intelligent Sensors, Sensor Networks and Information Processing (ISSNIP'05), Melbourne, Australia, 2005.
- [13] D. Tulone, S. Madden, PAQ: time series forecasting for approximate query answering in sensor networks, in: 3rd European Workshop on Wireless Sensor Networks (EWSN'06), Zurich, Switzerland, 2006.
- [14] S. Santini, K. Römer, An adaptive strategy for quality-based data reduction in wireless sensor networks, in: Proceedings of the 3rd International Conference on Networked Sensing Systems (INSS'06), Chicago, IL, USA, 2006.
- [15] Y. Xu, J. Winter, W.-C. Lee, Dual prediction-based reporting for object tracking sensor networks, in: Proceedings of the 1st International Conference on Mobile and Ubiquitous Systems: Networking and Services (MobiQutous'04), IEEE Computer Society, Los Alamitos, CA, USA, 2004, pp. 154–163.
- [16] Sensor Network in a Vineyard, GoodFood EU Integrated Project: Food Safety and Quality Monitoring with Microsystems, Project Website: (www3.unifi.it/midra/goodfood/).
- [17] G. Box, G. Jenkins, *Time Series Analysis: Forecasting and Control*, Holden-Day Series in Time Series Analysis, Holden-Day, San Francisco, CA, USA, 1976.
- [18] S. Makridakis, S. Wheelwright, R. Hyndman, *Forecasting: Methods and Applications*, third ed., Wiley, New York, 1998.
- [19] O. Maron, A.W. Moore, The racing algorithm: model selection for lazy learners, *Artificial Intell. Rev.* 11 (1–5) (1997) 193–225.
- [20] Intel Research Laboratories Berkeley: Intel Lab Data, Project website: (berkeley.intel-research.net/labdata/).
- [21] E.S. Gardner, Exponential smoothing: the state of the art, *J. Forecast.* 4 (1985) 1–38.
- [22] J.D.D. Gooijer, R. Hyndman, 25 Years of time series forecasting, *Int. J. Forecast.* 22 (3) (2006) 442–473.
- [23] D. Montgomery, L. Johnson, J. Gardiner, *Forecasting and Time Series Analysis*, second ed., McGraw-Hill, New York, 1990.
- [24] J. Hamilton, J. Hamilton, *Time Series Analysis*, Princeton University Press, Princeton, NJ, 1994.
- [25] W. Hoeffding, Probability inequalities for sums of bounded random variables, *J. Am. Statist. Assoc.* 58 (301) (1963) 13–30.
- [26] S. Alexander, *Adaptive Signal Processing: Theory and Applications*, Springer, New York, NY, USA, 1986.
- [27] A. Stenman, F. Gustafsson, L. Ljung, Just in time models for dynamical systems, in: Proceedings of the 35th IEEE Conference on Decision and Control, vol. 1, Kobe, Japan, 1996, pp. 1115–1120.
- [28] National Oceanic Atmospheric Administration's National Data Buoy Center, Historical data repository: (www.ndbc.noaa.gov/historical_data.shtml).
- [29] V. Raghunathan, C. Srivastava, Energy-aware wireless microsensor networks, *Signal Process. Mag. IEEE* 19 (2) (2002) 40–50.
- [30] M. Markou, S. Singh, Novelty detection: a review, *Signal Processing* 83 (12) (2003) 2481–2521.